

```
In [1]: # Newton method of sqrt
# this is actually Heron's method
# based on http://lomont.org/papers/2003/InvSqrt.pdf
# original guess within 3% of target inv_sqrt in that paper
# 0.1% after one newton step

###
# Here i show my init_guess for sqrt is within 12-13% of target sqrt
# 0.6-0.7% after one Newton step

# it's still much slower than default math.sqrt()
```

```
In [2]: import math
import random
```

```
In [3]: def f(x, I):
# function to use the newton method with
return x*x - I
```

```
In [4]: def f_p(x):
# the derivative of the f function (doesn't include I in this case, but in
general might)
return 2*x
```

```
In [5]: def init_sqrt_guess(x):
# returns a guess for sqrt(x)
# get the mantissa and exponent
m, e = math.frexp(x)
# we're looking for a new m and e such that (m*2**e)**2 ~= x
# for the exponent it's just half
e = e/2.0
# for the mantissa which is in [0.5,1[ (because x is positive), the best v
alue is obviously sqrt(m)
# we approximate that by sqrt(x) ~= x+k
# k = solve ((2 1^(3/2))/3 - 1^2/2 ) - ((2 0.5^(3/2))/3 - 0.5^2/2)-k/2 ==0
# I believe this k makes the integral of sqrt(x)-(x+k) in between 0.5,1 to
be zero
# and that this is optimal as a first guess (but doesn't necessarily optim
ized for the first step, the second step...)
# approximately
k = 0.111928812542301634
m = m+k
# init guess is
g = m*(2**e)
return g
```

```

In [6]: def newton_sqrt(I, max_steps=1, min_error=1e-200, add_noise=False, verbose=False):
    # computes the sqrt of I using the Newton method
    # default max_step=1 to mirror fast inverse sqrt paper
    # converges very fast so a very small min_error is possible
    # after 3 steps, 0% error

    I = float(I)

    # initial guess:
    x = init_sqrt_guess(I)

    step=0
    abs_error = abs(x*x-I)
    rel_error = 100.0*abs_error/I

    if verbose:
        print("Init guess: {}".format(x))
        print("Absolute Error: {}".format(abs_error))
        print("Relative Error: {:.10f} %".format(rel_error))

    # takes steps until max_steps or min_error reached, whichever comes first
    while (abs_error>min_error and step<max_steps):

        # add noise to prevent getting stuck?
        if add_noise: x += random.random()*min_error*2

        # Newton step
        x = x - f(x,I)/f_p(x)

        abs_error = abs(x*x-I)
        rel_error = 100.0*abs_error/I
        step += 1

        if verbose:
            print("After step {}".format(step))
            print("Current x: {}".format(x))
            print("Absolute Error: {}".format(abs_error))
            print("Relative Error: {:.10f} %".format(rel_error))

    return x

```

```
In [7]: def gen_rand_positive_32_bit_float():
        # hopefully a rand value over all the positive 32 bit floats?

        exponent_bits = [random.choice([0,1]) for _ in xrange(8)]
        mantissa_bits = [random.choice([0,1]) for _ in xrange(23)]

        # exponent, random int between -127 and 127
        exponent = 0
        for e,b in enumerate(exponent_bits[1:]):
            exponent += b*(2**e)
        # first bit defines exponent sign
        exponent *= (2*exponent_bits[0]-1)

        # mantissa
        mantissa = 1
        for e,b in enumerate(mantissa_bits):
            mantissa += b*(2**(-e))

        # print(mantissa_bits)
        # print(exponent_bits)
        # print(mantissa, exponent)

        # always positive
        r = mantissa*(2**exponent)

        return r
```

```
In [21]: # test some known values
        num_to_test = 100
        print(newton_sqrt(num_to_test))

10.0005472261
```

```
In [9]: # see the computation
        print(newton_sqrt(num_to_test, verbose=True))

Init guess: 6.92318420723
Absolute Error: 16.0695204327
Relative Error: 25.1086256761 %
After step 1
Current x: 8.08374269822
Absolute Error: 1.346896011
Relative Error: 2.1045250172 %
8.08374269822
```

```
In [10]: # test some random values
num_to_test = gen_rand_positive_32_bit_float()
print("Number to test: {}".format(num_to_test))

# my computation
print("My computation")
my_result = newton_sqrt(num_to_test, max_steps=6, min_error=1e-200, add_noise=False, verbose=True)

# python's computation
python_result = math.sqrt(num_to_test)

# compare
print("")
print("my results: {:.100f}".format(my_result))
print("python's   : {:.100f}".format(python_result))
my_error = abs(my_result*my_result-num_to_test)
print("Newton Absolute Error: {}".format(my_error))
print("Newton Relative Error: {:.5f}%".format(100.0*my_error/num_to_test))

py_error = abs(python_result*python_result-num_to_test)
print("Python Absolute Error: {}".format(py_error))
print("Python Relative Error: {:.5f}%".format(100.0*py_error/num_to_test))
```

Number to test: 3.96595994124e-24
 My computation
 Init guess: 1.82964175434e-12
 Absolute Error: 6.18370992004e-25
 Relative Error: 15.5919626311 %
 After step 1
 Current x: 1.99862865862e-12
 Absolute Error: 2.8556573817e-26
 Relative Error: 0.7200419127 %
 After step 2
 Current x: 1.99148461671e-12
 Absolute Error: 5.10373348747e-29
 Relative Error: 0.0012868848 %
 After step 3
 Current x: 1.99147180282e-12
 Absolute Error: 1.64195989659e-34
 Relative Error: 0.0000000041 %
 After step 4
 Current x: 1.99147180277e-12
 Absolute Error: 7.34683969264e-40
 Relative Error: 0.0000000000 %
 After step 5
 Current x: 1.99147180277e-12
 Absolute Error: 7.34683969264e-40
 Relative Error: 0.0000000000 %
 After step 6
 Current x: 1.99147180277e-12
 Absolute Error: 7.34683969264e-40
 Relative Error: 0.0000000000 %

my results: 0.000000000001991471802774202644131816104866208167172708654035773
 1792028062045574188232421875000000000
 python's : 0.000000000001991471802774203048028599578024252537977734078822322
 7718845009803771972656250000000000000
 Newton Absolute Error: 7.34683969264e-40
 Newton Relative Error: 0.00000%
 Python Absolute Error: 7.34683969264e-40
 Python Relative Error: 0.00000%

In [11]: `%%timeit -n 100000 -r 5
math.sqrt(num_to_test)`

100000 loops, best of 5: 91.2 ns per loop

In [12]: `%%timeit -n 100000 -r 5
newton_sqrt(num_to_test, max_steps=0)`

100000 loops, best of 5: 1.11 µs per loop

In [13]: `%%timeit -n 100000 -r 5
newton_sqrt(num_to_test, max_steps=1)`

100000 loops, best of 5: 1.73 µs per loop

```
In [14]: %%timeit -n 100000 -r 5
newton_sqrt(num_to_test, max_steps=2)
```

100000 loops, best of 5: 2.8 μ s per loop

```
In [15]: %%timeit -n 100000 -r 5
newton_sqrt(num_to_test, max_steps=3)
```

100000 loops, best of 5: 3.31 μ s per loop

```
In [16]: %%timeit -n 100000 -r 5
newton_sqrt(num_to_test, max_steps=4)
```

100000 loops, best of 5: 4.39 μ s per loop

```
In [17]: %%timeit -n 100000 -r 5
newton_sqrt(num_to_test, max_steps=5)
```

100000 loops, best of 5: 5.59 μ s per loop

```
In [18]: # seems about 20 times slower with maxstep=1
```

In [19]: *# test over a bunch of floats, to evaluate the quality of my init_guess*

```
num_samples = 1000000

for max_steps in range(6):
    print("Testing with max_steps={}".format(max_steps))

    max_abs_error = -1
    max_abs_err_num = -1
    max_abs_err_res = -1

    max_rel_error = -1
    max_rel_err_num = -1
    max_rel_err_res = -1

    max_sqrt_abs_error = -1
    max_sqrt_abs_err_num = -1
    max_sqrt_abs_err_res = -1

    max_sqrt_rel_error = -1
    max_sqrt_abs_err_num = -1
    max_sqrt_abs_err_res = -1

    average_abs_err = 0
    average_rel_err = 0

    for _ in xrange(num_samples):
        # generate a new random value
        num_to_test = gen_rand_positive_32_bit_float()

        # compute my sqrt
        my_result = newton_sqrt(num_to_test, max_steps=max_steps, min_error=1e
-100, add_noise=False, verbose=False)

        my_abs_error = abs(my_result*my_result-num_to_test)
        my_rel_error = 100.0*my_abs_error/num_to_test

        sqrt_abs_error = abs(my_result-math.sqrt(num_to_test))
        sqrt_rel_error = 100.0*sqrt_abs_error/math.sqrt(num_to_test)

        average_abs_err += my_abs_error
        average_rel_err += my_rel_error

        if (max_abs_error == -1) or (my_abs_error>max_abs_error):
            max_abs_error = my_abs_error
            max_abs_err_num = num_to_test
            max_abs_err_res = my_result*my_result

        if (max_rel_error == -1) or (my_rel_error>max_rel_error):
            max_rel_error = my_rel_error
            max_rel_err_num = num_to_test
            max_rel_err_res = my_result*my_result

        if (max_sqrt_abs_error == -1) or (sqrt_abs_error>max_sqrt_abs_error):
            max_sqrt_abs_error = sqrt_abs_error
            max_sqrt_abs_err_num = math.sqrt(num_to_test)
```

```

        max_sqrt_abs_err_res = my_result

    if (max_sqrt_rel_error == -1) or (sqrt_rel_error > max_sqrt_rel_error):
        max_sqrt_rel_error = sqrt_rel_error
        max_sqrt_rel_err_num = math.sqrt(num_to_test)
        max_sqrt_rel_err_res = my_result

    average_abs_err /= num_samples
    average_rel_err /= num_samples

    print("Max absolute error: {} (result: {} target: {})".format(max_abs_err
or, max_abs_err_res, max_abs_err_num))
    print("Max relative error: {}% (result: {} target: {})".format(max_rel_er
ror, max_rel_err_res, max_rel_err_num))
    print("Max sqrt absolute error: {} (result: {} target: {})".format(max_sq
rt_abs_error, max_sqrt_abs_err_res, max_sqrt_abs_err_num))
    print("Max sqrt relative error: {}% (result: {} target: {})".format(max_s
qrt_rel_error, max_sqrt_rel_err_res, max_sqrt_rel_err_num))
    print("Average absolute error: {}".format(average_abs_err))
    print("Average relative error: {}%".format(average_rel_err))
    print('')

```


Testing with max_steps=0

Max absolute error: 8.54318421654e+37 (result: 2.54887966084e+38 target: 3.40319808249e+38)

Max relative error: 25.1085237642% (result: 4.25709668749e-14 target: 5.68435408336e-14)

Max sqrt absolute error: 2.48254777573e+18 (result: 1.59652111193e+19 target: 1.8447758895e+19)

Max sqrt relative error: 13.4601385281% (result: 2.06327329443e-07 target: 2.38418834897e-07)

Average absolute error: 3.2547308125e+35

Average relative error: 12.6443008656%

Testing with max_steps=1

Max absolute error: 7.15771615647e+36 (result: 3.47490302324e+38 target: 3.40332586167e+38)

Max relative error: 2.10452501718% (result: 0.510522625086 target: 0.5)

Max sqrt absolute error: 1.92986549665e+17 (result: 1.86410917686e+19 target: 1.84481052189e+19)

Max sqrt relative error: 1.04678372773% (result: 0.71450865991 target: 0.707106781187)

Average absolute error: 1.51072690122e+34

Average relative error: 0.605191753169%

Testing with max_steps=2

Max absolute error: 3.68135084371e+34 (result: 3.40423918793e+38 target: 3.40387105284e+38)

Max relative error: 0.0108442515549% (result: 64.0070013628 target: 64.0000610352)

Max sqrt absolute error: 9.97651641741e+14 (result: 1.8450580446e+19 target: 1.84495827943e+19)

Max sqrt relative error: 0.00542197878817% (result: 8.00043757321 target: 8.0000038147)

Average absolute error: 3.9841555815e+31

Average relative error: 0.0017124450769%

Testing with max_steps=3

Max absolute error: 9.96085135454e+29 (result: 3.40369906797e+38 target: 3.40369905801e+38)

Max relative error: 2.9396692684e-07% (result: 17179871282.5 target: 17179871232.0)

Max sqrt absolute error: 26995470336.0 (result: 1.84491166942e+19 target: 1.84491166672e+19)

Max sqrt relative error: 1.46983461079e-07% (result: 131072.008005 target: 131072.007812)

Average absolute error: 5.53729386455e+26

Average relative error: 2.36442946349e-08%

Testing with max_steps=4

Max absolute error: 1.51115727452e+23 (result: 4.32647040487e+38 target: 4.32647040487e+38)

Max relative error: 4.42161718457e-14% (result: 4.52323192288e+15 target: 4.52323192288e+15)

Max sqrt absolute error: 4096.0 (result: 2.07109436902e+19 target: 2.07109436902e+19)

Max sqrt relative error: 2.22042275621e-14% (result: 3.05178982655e-05 target: 3.05178982655e-05)

Average absolute error: 2.16277722264e+20

Average relative error: 8.15315896066e-15%

Testing with max_steps=5

Max absolute error: 7.55578637259e+22 (result: 4.8917930835e+38 target: 4.8917930835e+38)

Max relative error: 2.22044393167e-14% (result: 2.35099094372e-38 target: 2.35099094372e-38)

Max sqrt absolute error: 4096.0 (result: 2.02234731734e+19 target: 2.02234731734e+19)

Max sqrt relative error: 2.22042725595e-14% (result: 0.062500528989 target: 0.062500528989)

Average absolute error: 2.10425059926e+20

Average relative error: 7.84093481805e-15%